

```

import pandas as pd
import numpy as np

# Data exploration
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from scipy.stats import skew

# Data wrangling and feature engineering
from sklearn import preprocessing
from sklearn.feature_selection import chi2
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from sklearn.preprocessing import LabelEncoder

# Deep learning
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, BatchNormalization, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import accuracy_score, confusion_matrix

df = pd.read_csv('/loan_data.csv')
df.head()

```

	credit.policy		purpose	int.rate	installment
log.annual.inc \					
0	1	debt_consolidation	0.1189	829.10	
11.350407					
1	1	credit_card	0.1071	228.22	
11.082143					
2	1	debt_consolidation	0.1357	366.86	
10.373491					
3	1	debt_consolidation	0.1008	162.34	
11.350407					
4	1	credit_card	0.1426	102.92	
11.299732					

	dti	fico	days.with.cr.line	revol.bal	revol.util
inq.last.6mths \					
0	19.48	737	5639.958333	28854	52.1
0					
1	14.29	707	2760.000000	33623	76.7
0					
2	11.63	682	4710.000000	3511	25.6

```

1
3    8.10    712          2699.958333    33667    73.2
1
4    14.97    667          4066.000000    4740    39.5
0

```

```

delinq.2yrs  pub.rec  not.fully.paid
0            0        0            0
1            0        0            0
2            0        0            0
3            0        0            0
4            1        0            0

```

```
df.shape
```

```
(9578, 14)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 9578 entries, 0 to 9577
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	credit.policy	9578 non-null	int64
1	purpose	9578 non-null	object
2	int.rate	9578 non-null	float64
3	installment	9578 non-null	float64
4	log.annual.inc	9578 non-null	float64
5	dti	9578 non-null	float64
6	fico	9578 non-null	int64
7	days.with.cr.line	9578 non-null	float64
8	revol.bal	9578 non-null	int64
9	revol.util	9578 non-null	float64
10	inq.last.6mths	9578 non-null	int64
11	delinq.2yrs	9578 non-null	int64
12	pub.rec	9578 non-null	int64
13	not.fully.paid	9578 non-null	int64

```
dtypes: float64(6), int64(7), object(1)
```

```
memory usage: 1.0+ MB
```

```
df['not.fully.paid'].value_counts()
```

```
0    8045
```

```
1    1533
```

```
Name: not.fully.paid, dtype: int64
```

```
not_fully_paid0 = df[df['not.fully.paid']==0]
```

```
not_fully_paid1 = df[df['not.fully.paid']==1]
```

```
print('not_fully_paid1      :',not_fully_paid1.shape,'\n',
      'not_fully_paid0      :',not_fully_paid0.shape)
```

```

not_fully_paid1    : (1533, 14)
not_fully_paid0    : (8045, 14)

from sklearn.utils import resample
df_minority = resample(not_fully_paid1,n_samples = 8045,replace=True)

df = pd.concat([not_fully_paid0,df_minority])

from sklearn.utils import shuffle
df = shuffle(df)

df['not.fully.paid'].value_counts()

1      8045
0      8045
Name: not.fully.paid, dtype: int64

# Separate data to include numerical data only
num_data = df[["int.rate", "installment", "log.annual.inc", "dti",
               "fico", "days.with.cr.line", "revol.bal",
               "revol.util", "not.fully.paid"]]

# Check the features in the categorical data
cat_data = df[["credit.policy", "purpose", "inq.last.6mths",
               "delinq.2yrs", "not.fully.paid"]]

```

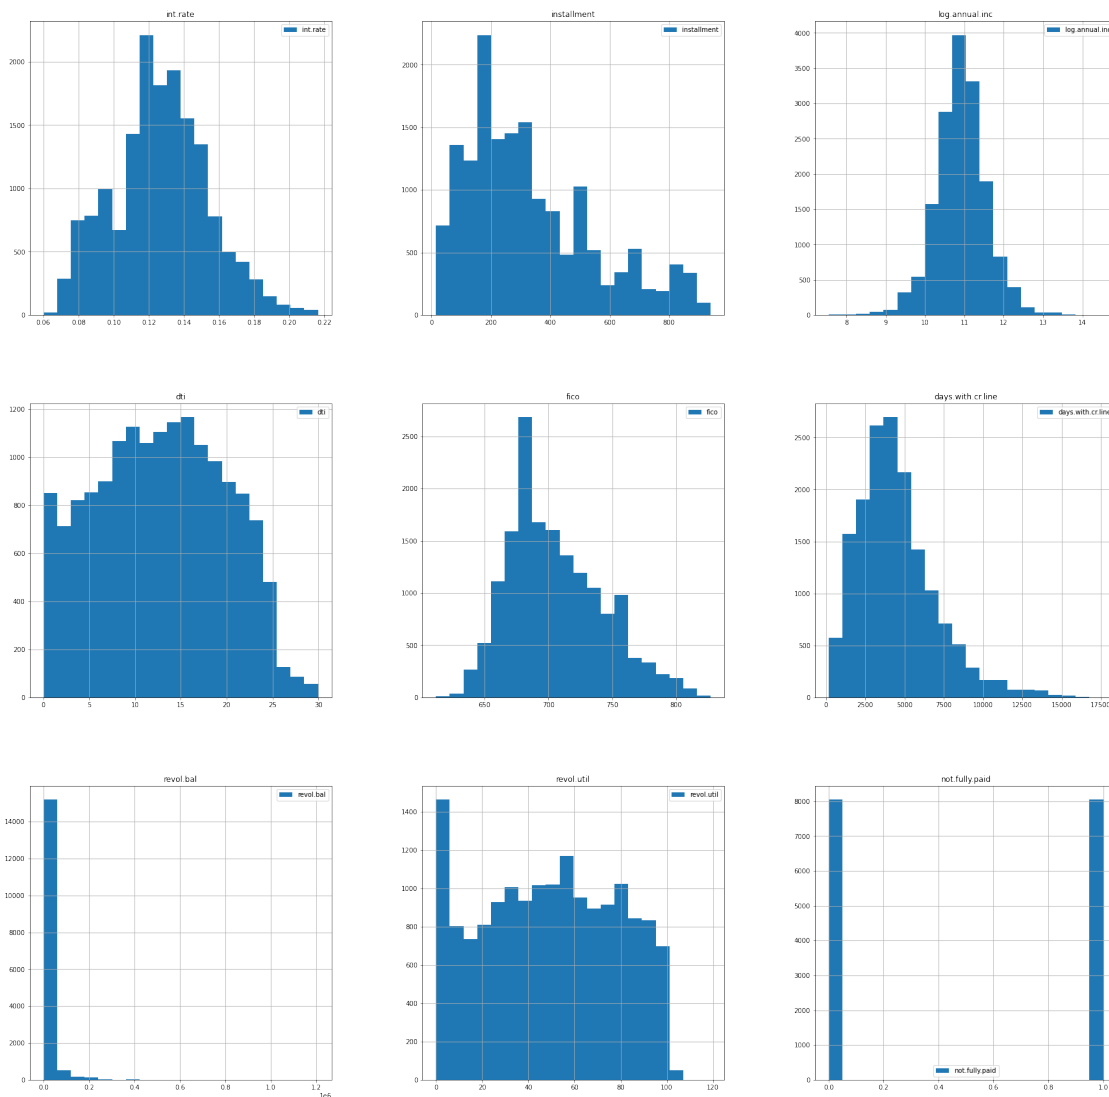
```
num_data.describe()
```

	int.rate	installment	log.annual.inc	dti
fico \				
count	16090.000000	16090.000000	16090.000000	16090.000000
mean	0.126717	331.233010	10.918552	12.779421
std	0.026932	216.046425	0.639732	6.968271
min	0.060000	15.670000	7.547502	0.000000
25%	0.110300	167.120000	10.524064	7.280000
50%	0.126100	277.450000	10.915088	12.890000
75%	0.143800	469.780000	11.289832	18.240000
max	0.216400	940.140000	14.528354	29.960000

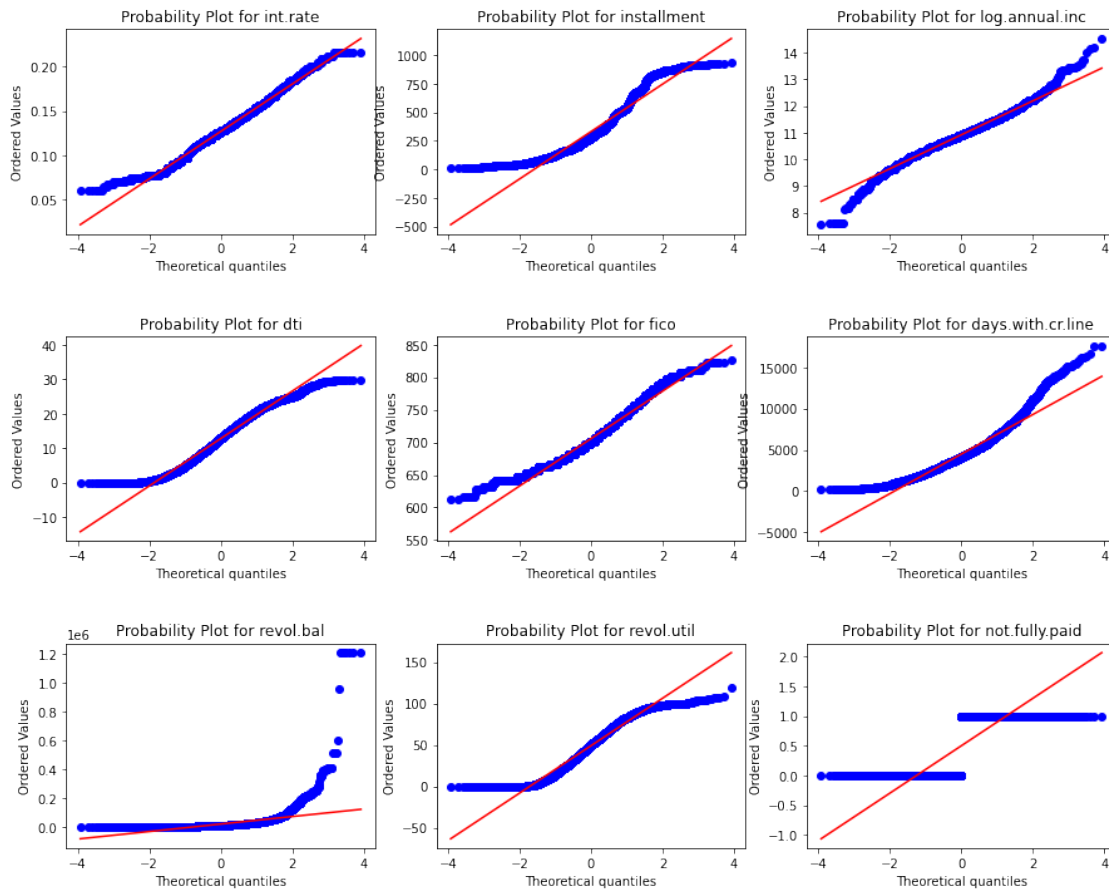
	days.with.cr.line	revol.bal	revol.util	not.fully.paid
count	16090.000000	1.609000e+04	16090.000000	16090.000000
mean	4494.871600	1.921108e+04	48.959622	0.500000
std	2474.362329	4.398412e+04	29.142480	0.500016

min	178.958333	0.000000e+00	0.000000	0.000000
25%	2790.000000	3.163250e+03	24.900000	0.000000
50%	4110.000000	8.764000e+03	49.800000	0.500000
75%	5670.041667	1.914400e+04	73.500000	1.000000
max	17639.958330	1.207359e+06	119.000000	1.000000

```
num_data.hist(figsize=(30,30),bins=20,legend=True)
plt.show()
```



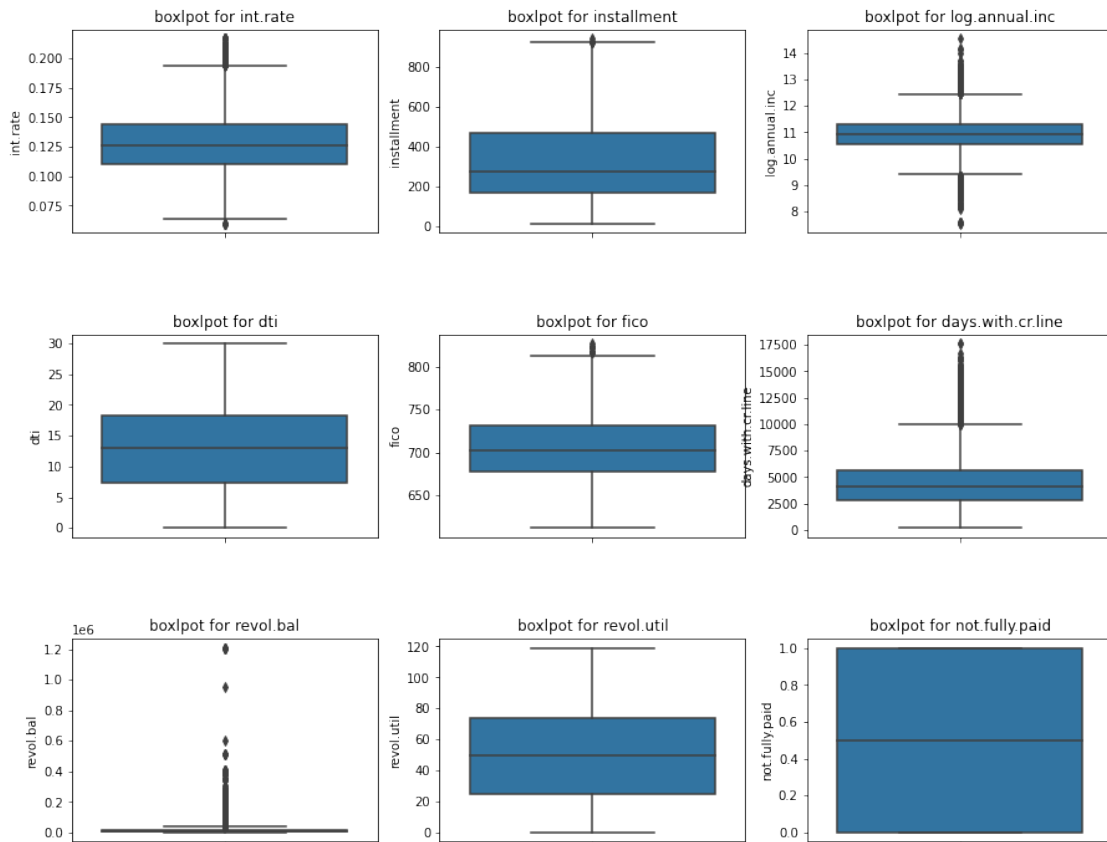
```
fig,ax= plt.subplots(3,3,figsize=(15,12),sharex=True)
fig.subplots_adjust(hspace=0.5)
for i,cols in enumerate(num_data):
    axx = plt.subplot(3,3,i+1)
    stats.probplot(num_data[cols],plot = axx)
    axx.set_title(f"Probability Plot for {cols}")
plt.show()
```



*# it can be seen that the variables such as revol.bal,
days.with.cr.line, installment, fico, and revol.util
they may contain outliers because the values in these variables do
not fall well around the best fit line.*

```
fig,ax = plt.subplots(3,3,figsize=(15,12),sharex=True)
fig.subplots_adjust(hspace=0.5)
```

```
for i,j in enumerate(num_data):
    ax = plt.subplot(3,3,i+1)
    sns.boxplot(y=num_data[j])
    ax.set_title(f'boxplot for {j}')
plt.show()
```



From the graphs above, it can be seen that the outliers exist in the variables

such as the following: `int.rate`, `installment`, `log.annual.inc`, `fico`, `days.with.cr.line` and `revol.bal`. These outliers will be handles later.

Converting categorical feature into numerical feature

```
cat_data = cat_data.copy()
le = preprocessing.LabelEncoder()
cat_data["purpose"] =
le.fit_transform(cat_data["purpose"].astype(str))
cat_data.head()
```

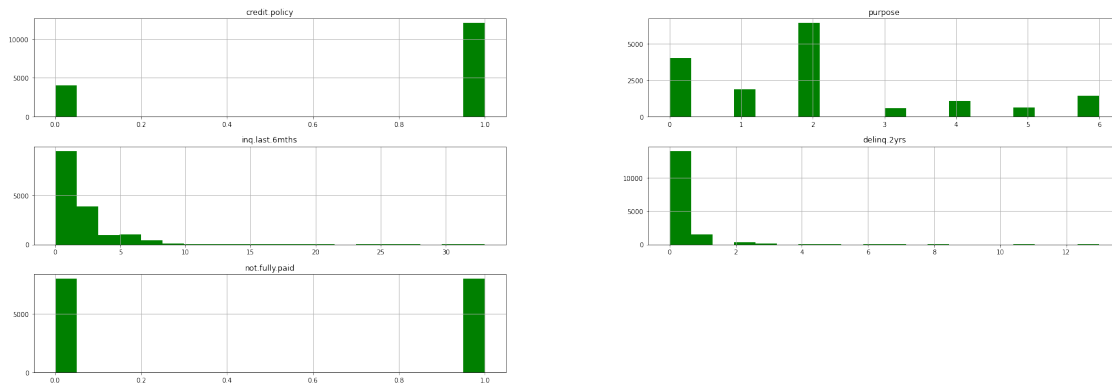
	credit.policy	purpose	inq.last.6mths	delinq.2yrs
not.fully.paid				
8008	0	6	5	2
1				
7050	1	2	0	0
0				
6386	1	2	5	0
0				
8744	0	1	8	1
1				
8357	0	0	5	1
1				

```
cat_data.describe()
```

	credit.policy	purpose	inq.last.6mths	delinq.2yrs	\
count	16090.000000	16090.000000	16090.000000	16090.000000	
mean	0.747296	2.029770	1.869733	0.170230	
std	0.434576	1.788958	2.521743	0.536373	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	1.000000	0.000000	0.000000	
50%	1.000000	2.000000	1.000000	0.000000	
75%	1.000000	2.000000	3.000000	0.000000	
max	1.000000	6.000000	33.000000	13.000000	

	not.fully.paid
count	16090.000000
mean	0.500000
std	0.500016
min	0.000000
25%	0.000000
50%	0.500000
75%	1.000000
max	1.000000

```
cat_data.hist(figsize=(30,10),bins=20,color='g')
plt.rcParams['font.size'] = '20'
plt.show()
```



It can be seen that most of the categorical data is positively skewed.

Most clients satisfied the credit policy.

Most clients decided to take the loan for purposes of loan consolidation

```
le = preprocessing.LabelEncoder()
df["purpose"] = le.fit_transform(df["purpose"].astype(str))
df.head()
```

	credit.policy	purpose	int.rate	installment	log.annual.inc
dti \					
8008	0	6	0.1280	503.97	11.002100

22.44					
7050	1	2	0.1311	398.19	11.238489
20.42					
6386	1	2	0.1148	478.05	11.277203
6.88					
8744	0	1	0.1064	244.27	10.596635
2.07					
8357	0	0	0.1324	169.05	10.239960
19.54					

	fico	days.with.cr.line	revol.bal	revol.util	
inq.last.6mths \					
8008	687	4859.041667	23	1.3	5
7050	692	4267.041667	12764	44.3	0
6386	752	7592.000000	18003	35.0	5
8744	702	11730.000000	599	3.2	8
8357	667	3960.041667	13542	56.4	5

	delinq.2yrs	pub.rec	not.fully.paid
8008	2	0	1
7050	0	1	0
6386	0	0	0
8744	1	1	1
8357	1	0	1

```
df.isnull().sum()
```

```
credit.policy      0
purpose            0
int.rate           0
installment        0
log.annual.inc     0
dti                0
fico               0
days.with.cr.line 0
revol.bal          0
revol.util         0
inq.last.6mths     0
delinq.2yrs        0
pub.rec            0
not.fully.paid     0
dtype: int64
```

```
num_data.describe()
```


	int.rate	installment	log.annual.inc	dti
fico \				
count	16090.000000	16090.000000	16090.000000	16090.000000
16090.000000				
mean	0.126717	331.233010	10.918552	12.779421
705.695152				
std	0.026932	216.046425	0.639732	6.968271
36.941948				
min	0.060000	15.670000	7.547502	0.000000
612.000000				
25%	0.110300	167.120000	10.524064	7.280000
677.000000				
50%	0.126100	277.450000	10.915088	12.890000
702.000000				
75%	0.143800	469.780000	11.289832	18.240000
732.000000				
max	0.216400	940.140000	14.528354	29.960000
827.000000				

	days.with.cr.line	revol.bal	revol.util	not.fully.paid
count	16090.000000	1.609000e+04	16090.000000	16090.000000
mean	4494.871600	1.921108e+04	48.959622	0.500000
std	2474.362329	4.398412e+04	29.142480	0.500016
min	178.958333	0.000000e+00	0.000000	0.000000
25%	2790.000000	3.163250e+03	24.900000	0.000000
50%	4110.000000	8.764000e+03	49.800000	0.500000
75%	5670.041667	1.914400e+04	73.500000	1.000000
max	17639.958330	1.207359e+06	119.000000	1.000000

Upper bounded outliers

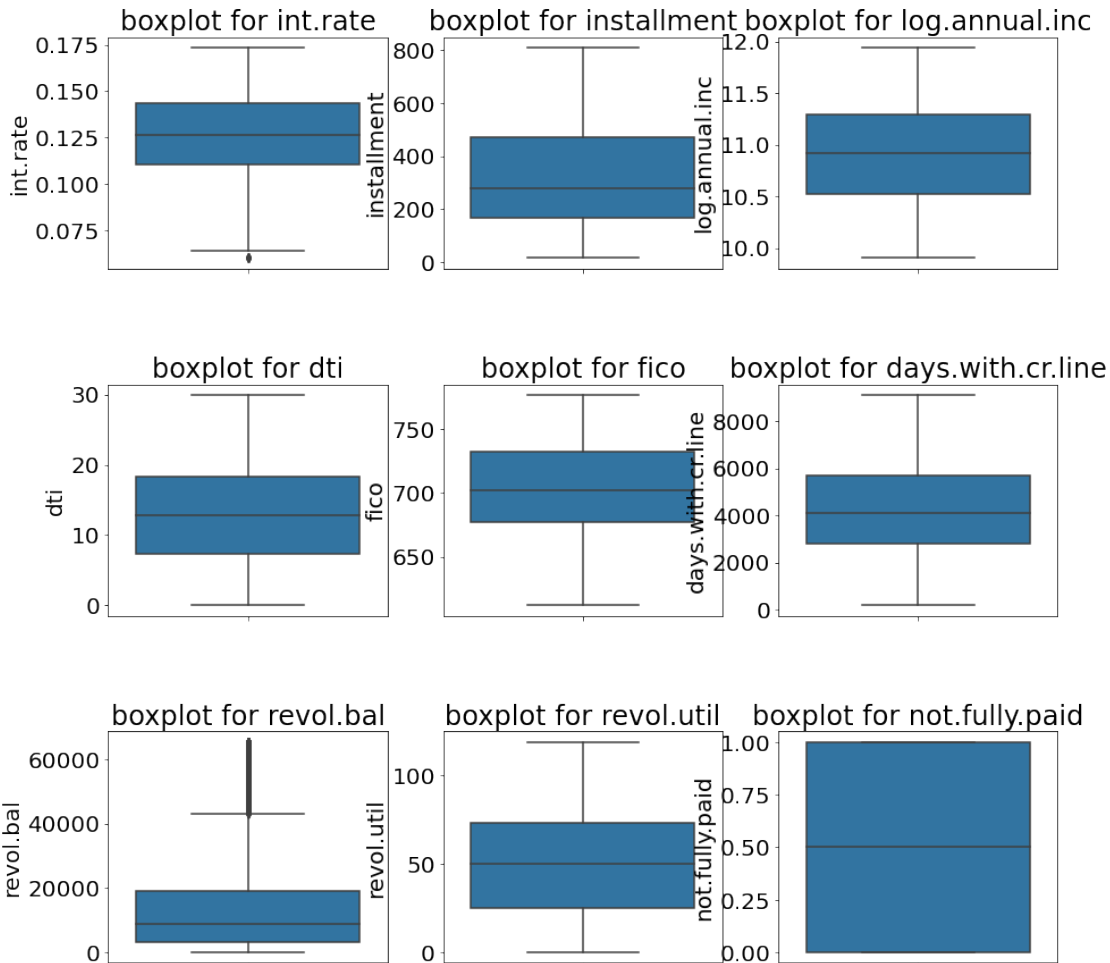
```
for var in ['int.rate', 'installment', 'fico', 'days.with.cr.line',
            'revol.bal', 'not.fully.paid']:
    df[var].clip(upper = df[var].quantile(0.95), inplace=True)
```

Lower and Upper bounded outliers

```
for var in ['log.annual.inc']:
    df[var] = df[var].clip(upper = df[var].quantile(0.95), lower =
df[var].quantile(0.05))
```

```
fig, ax = plt.subplots(3, 3, figsize=(15, 15))
plt.subplots_adjust(hspace=0.5)
```

```
for i, j in enumerate(df[num_data.columns]):
    ax = plt.subplot(3, 3, i+1)
    sns.boxplot(y=df[j])
    ax.set_title(f'boxplot for {j}')
plt.show()
```



Check for skewness in the numerical features

```
vars_skewed = df[num_data.columns].apply(lambda x:
skew(x)).sort_values(ascending=False)
vars_skewed
```

```
revol.bal          1.709400
installment        0.781616
days.with.cr.line 0.470738
fico               0.362407
dti               -0.017470
log.annual.inc     0.006915
not.fully.paid     0.000000
revol.util        -0.034960
int.rate          -0.098591
dtype: float64
```

```
vars_skewed = vars_skewed[vars_skewed>0.3]
vars_skewed
```

```
revol.bal          1.709400
installment        0.781616
days.with.cr.line 0.470738
```

```
fico          0.362407
dtype: float64
```

```
for i in vars_skewed.index:
    df[i] = np.log1p(df[i])
```

```
vars_skewed= df[num_data.columns].apply(lambda x:
skew(x)).sort_values(ascending=False)
vars_skewed
```

```
fico          0.273172
dti           0.017470
log.annual.inc 0.006915
not.fully.paid 0.000000
revol.util    -0.034960
int.rate      -0.098591
installment   -0.586084
days.with.cr.line -1.146639
revol.bal     -2.284327
dtype: float64
```

```
# for i in cat_data.columns:
#     print(df[i].value_counts(),'\n')
```

```
x_num = df[num_data.columns.drop('not.fully.paid')]
x_num
```

	int.rate	installment	log.annual.inc	dti	fico	\
8008	0.1280	6.224499	11.002100	22.44	6.533789	
7050	0.1311	5.989437	11.238489	20.42	6.541030	
6386	0.1148	6.171805	11.277203	6.88	6.624065	
8744	0.1064	5.502360	10.596635	2.07	6.555357	
8357	0.1324	5.136093	10.239960	19.54	6.504288	
...	
1388	0.1209	6.396329	11.112448	19.22	6.555357	
5621	0.1287	5.376481	11.167261	11.09	6.533789	
2714	0.1347	6.011463	10.149097	9.20	6.583409	
3009	0.1347	5.607235	10.550931	12.59	6.533789	
8697	0.1576	4.173618	9.903688	25.07	6.504288	


	days.with.cr.line	revol.bal	revol.util
8008	8.488802	3.178054	1.3
7050	8.358910	9.454462	44.3
6386	8.934982	9.798349	35.0
8744	9.118333	6.396930	3.2
8357	8.284262	9.513625	56.4
...
1388	8.560444	10.038805	64.5
5621	8.804025	9.414750	40.0

2714	9.118333	9.339437	94.8
3009	8.137968	8.620832	60.9
8697	7.186113	8.596374	35.6

[16090 rows x 8 columns]

```
y = df[['not.fully.paid']]
```

```
plt.figure(figsize=(40,15))
sns.heatmap(x_num.corr(),annot=True,cmap='bwr')
plt.show()
```



```
matrix = x_num.corr().unstack()
sorted_pairs = matrix.sort_values()

sorted_pairs[abs(sorted_pairs)>0.7]
```

int.rate	int.rate	1.0
days.with.cr.line	days.with.cr.line	1.0
fico	fico	1.0
dti	dti	1.0
log.annual.inc	log.annual.inc	1.0
installment	installment	1.0
revol.bal	revol.bal	1.0
revol.util	revol.util	1.0

dtype: float64

```
from sklearn.svm import SVR
estimator = SVR(kernel='linear')
rfe = RFE(estimator,n_features_to_select=5,step = 1)
rfe= rfe.fit(x_num,y.values)
```

/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993
: DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993
: DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993
: DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)

list(zip(x_num.columns , rfe.support_, rfe.ranking_))

[('int.rate', True, 1),
 ('installment', True, 1),
 ('log.annual.inc', True, 1),
 ('dti', False, 3),
 ('fico', True, 1),
 ('days.with.cr.line', True, 1),
 ('revol.bal', False, 2),
 ('revol.util', False, 4)]

selected_cols = x_num.columns[rfe.support_]
selected_cols

Index(['int.rate', 'installment', 'log.annual.inc', 'fico',
      'days.with.cr.line'],
      dtype='object')

best_num_cols = df[selected_cols]

# select best categorical columns

cat_vars = df[cat_data.columns].drop('not.fully.paid', axis=1)
cat_vars.head()

   credit.policy  purpose  inq.last.6mths  delinq.2yrs
8008           0         6              5           2
7050           1         2              0           0
6386           1         2              5           0
8744           0         1              8           1
8357           0         0              5           1

f_p_values = chi2(cat_vars, df['not.fully.paid'])
f_p_values

```

```
(array([ 155.64071856,  116.07241495, 1634.36191996,    7.05403432]),
 array([1.01432920e-35, 4.58254427e-27, 0.00000000e+00, 7.90869323e-
03]))
```

wseeing what value belongs to which feature

```
pd.Series(f_p_values[1],cat_vars.columns).sort_values(ascending=True)
```

```
inq.last.6mths    0.000000e+00
credit.policy     1.014329e-35
purpose           4.582544e-27
delinq.2yrs       7.908693e-03
dtype: float64
```

all four categorical columns have p value less than 0.05 . so consider all of them

```
x = df[['int.rate', 'installment', 'log.annual.inc', 'dti', 'fico',
'inq.last.6mths','credit.policy', 'purpose','delinq.2yrs']]
x.head()
```

	int.rate	installment	log.annual.inc	dti	fico
inq.last.6mths \					
8008	0.1280	6.224499	11.002100	22.44	6.533789
5					
7050	0.1311	5.989437	11.238489	20.42	6.541030
0					
6386	0.1148	6.171805	11.277203	6.88	6.624065
5					
8744	0.1064	5.502360	10.596635	2.07	6.555357
8					
8357	0.1324	5.136093	10.239960	19.54	6.504288
5					

	credit.policy	purpose	delinq.2yrs
8008	0	6	2
7050	1	2	0
6386	1	2	0
8744	0	1	1
8357	0	0	1

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size =
0.2, random_state = 42)
```

Scale the data

```
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
model = keras.Sequential(
[
keras.layers.Dense(
256, activation="relu", input_shape=[9]),
```

```

        keras.layers.Dense(256, activation="relu"),
        keras.layers.Dropout(0.3),
        keras.layers.Dense(256, activation="relu"),
        keras.layers.Dropout(0.3),
        keras.layers.Dense(1, activation="sigmoid"),
    ]
)
model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 256)	2560
dense_5 (Dense)	(None, 256)	65792
dropout_2 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 1)	257
Total params: 134,401		
Trainable params: 134,401		
Non-trainable params: 0		

```

model.compile(optimizer = 'Adam', loss = 'binary_crossentropy',
metrics = ['binary_accuracy'])

```

```

early_stopping = keras.callbacks.EarlyStopping(patience=10,
min_delta=0.001, restore_best_weights=True)

```

```

history = model.fit(
    x_train, y_train,
    validation_data=(x_test, y_test),
    batch_size=256,
    epochs=100,
    callbacks=[early_stopping],
    verbose=1,
)

```

Epoch 1/100

51/51 [=====] - 1s 21ms/step - loss: 0.4797 - binary_accuracy: 0.7547 - val_loss: 0.5214 - val_binary_accuracy: 0.7390

Epoch 2/100

51/51 [=====] - 1s 14ms/step - loss: 0.4791 -

binary_accuracy: 0.7571 - val_loss: 0.5165 - val_binary_accuracy:
0.7362
Epoch 3/100
51/51 [=====] - 1s 13ms/step - loss: 0.4688 -
binary_accuracy: 0.7609 - val_loss: 0.5187 - val_binary_accuracy:
0.7523
Epoch 4/100
51/51 [=====] - 1s 12ms/step - loss: 0.4743 -
binary_accuracy: 0.7557 - val_loss: 0.5071 - val_binary_accuracy:
0.7523
Epoch 5/100
51/51 [=====] - 1s 12ms/step - loss: 0.4654 -
binary_accuracy: 0.7641 - val_loss: 0.5140 - val_binary_accuracy:
0.7464
Epoch 6/100
51/51 [=====] - 1s 13ms/step - loss: 0.4545 -
binary_accuracy: 0.7702 - val_loss: 0.5073 - val_binary_accuracy:
0.7483
Epoch 7/100
51/51 [=====] - 1s 13ms/step - loss: 0.4581 -
binary_accuracy: 0.7651 - val_loss: 0.5123 - val_binary_accuracy:
0.7495
Epoch 8/100
51/51 [=====] - 1s 13ms/step - loss: 0.4529 -
binary_accuracy: 0.7777 - val_loss: 0.4994 - val_binary_accuracy:
0.7604
Epoch 9/100
51/51 [=====] - 1s 13ms/step - loss: 0.4440 -
binary_accuracy: 0.7777 - val_loss: 0.5053 - val_binary_accuracy:
0.7592
Epoch 10/100
51/51 [=====] - 1s 13ms/step - loss: 0.4475 -
binary_accuracy: 0.7770 - val_loss: 0.4913 - val_binary_accuracy:
0.7548
Epoch 11/100
51/51 [=====] - 1s 13ms/step - loss: 0.4470 -
binary_accuracy: 0.7822 - val_loss: 0.5132 - val_binary_accuracy:
0.7557
Epoch 12/100
51/51 [=====] - 1s 12ms/step - loss: 0.4378 -
binary_accuracy: 0.7825 - val_loss: 0.5008 - val_binary_accuracy:
0.7632
Epoch 13/100
51/51 [=====] - 1s 19ms/step - loss: 0.4309 -
binary_accuracy: 0.7885 - val_loss: 0.5036 - val_binary_accuracy:
0.7601
Epoch 14/100
51/51 [=====] - 2s 37ms/step - loss: 0.4396 -
binary_accuracy: 0.7804 - val_loss: 0.4956 - val_binary_accuracy:
0.7713

Epoch 15/100
51/51 [=====] - 1s 22ms/step - loss: 0.4283 -
binary_accuracy: 0.7894 - val_loss: 0.4867 - val_binary_accuracy:
0.7623
Epoch 16/100
51/51 [=====] - 1s 25ms/step - loss: 0.4269 -
binary_accuracy: 0.7920 - val_loss: 0.4931 - val_binary_accuracy:
0.7707
Epoch 17/100
51/51 [=====] - 2s 30ms/step - loss: 0.4222 -
binary_accuracy: 0.7930 - val_loss: 0.4842 - val_binary_accuracy:
0.7672
Epoch 18/100
51/51 [=====] - 1s 20ms/step - loss: 0.4153 -
binary_accuracy: 0.7958 - val_loss: 0.5007 - val_binary_accuracy:
0.7784
Epoch 19/100
51/51 [=====] - 1s 15ms/step - loss: 0.4152 -
binary_accuracy: 0.7960 - val_loss: 0.4847 - val_binary_accuracy:
0.7787
Epoch 20/100
51/51 [=====] - 1s 16ms/step - loss: 0.4123 -
binary_accuracy: 0.7997 - val_loss: 0.4902 - val_binary_accuracy:
0.7660
Epoch 21/100
51/51 [=====] - 1s 27ms/step - loss: 0.4096 -
binary_accuracy: 0.7995 - val_loss: 0.4814 - val_binary_accuracy:
0.7775
Epoch 22/100
51/51 [=====] - 1s 28ms/step - loss: 0.4090 -
binary_accuracy: 0.8025 - val_loss: 0.4872 - val_binary_accuracy:
0.7759
Epoch 23/100
51/51 [=====] - 1s 28ms/step - loss: 0.4083 -
binary_accuracy: 0.7990 - val_loss: 0.4944 - val_binary_accuracy:
0.7797
Epoch 24/100
51/51 [=====] - 2s 42ms/step - loss: 0.4048 -
binary_accuracy: 0.8034 - val_loss: 0.4883 - val_binary_accuracy:
0.7747
Epoch 25/100
51/51 [=====] - 1s 28ms/step - loss: 0.3976 -
binary_accuracy: 0.8083 - val_loss: 0.4966 - val_binary_accuracy:
0.7800
Epoch 26/100
51/51 [=====] - 1s 15ms/step - loss: 0.3977 -
binary_accuracy: 0.8082 - val_loss: 0.4838 - val_binary_accuracy:
0.7871
Epoch 27/100
51/51 [=====] - 1s 20ms/step - loss: 0.3960 -

binary_accuracy: 0.8063 - val_loss: 0.4694 - val_binary_accuracy:
0.7915
Epoch 28/100
51/51 [=====] - 1s 18ms/step - loss: 0.3939 -
binary_accuracy: 0.8117 - val_loss: 0.4900 - val_binary_accuracy:
0.7778
Epoch 29/100
51/51 [=====] - 1s 24ms/step - loss: 0.3916 -
binary_accuracy: 0.8101 - val_loss: 0.4831 - val_binary_accuracy:
0.7843
Epoch 30/100
51/51 [=====] - 1s 17ms/step - loss: 0.3858 -
binary_accuracy: 0.8159 - val_loss: 0.4820 - val_binary_accuracy:
0.7915
Epoch 31/100
51/51 [=====] - 1s 17ms/step - loss: 0.3834 -
binary_accuracy: 0.8188 - val_loss: 0.4752 - val_binary_accuracy:
0.7955
Epoch 32/100
51/51 [=====] - 1s 18ms/step - loss: 0.3818 -
binary_accuracy: 0.8167 - val_loss: 0.4830 - val_binary_accuracy:
0.7819
Epoch 33/100
51/51 [=====] - 1s 15ms/step - loss: 0.3853 -
binary_accuracy: 0.8181 - val_loss: 0.4826 - val_binary_accuracy:
0.7899
Epoch 34/100
51/51 [=====] - 1s 16ms/step - loss: 0.3746 -
binary_accuracy: 0.8252 - val_loss: 0.4737 - val_binary_accuracy:
0.7884
Epoch 35/100
51/51 [=====] - 1s 18ms/step - loss: 0.3761 -
binary_accuracy: 0.8222 - val_loss: 0.4975 - val_binary_accuracy:
0.7825
Epoch 36/100
51/51 [=====] - 1s 17ms/step - loss: 0.3829 -
binary_accuracy: 0.8176 - val_loss: 0.4798 - val_binary_accuracy:
0.7946
Epoch 37/100
51/51 [=====] - 1s 23ms/step - loss: 0.3764 -
binary_accuracy: 0.8198 - val_loss: 0.4674 - val_binary_accuracy:
0.7993
Epoch 38/100
51/51 [=====] - 1s 27ms/step - loss: 0.3679 -
binary_accuracy: 0.8276 - val_loss: 0.4702 - val_binary_accuracy:
0.7965
Epoch 39/100
51/51 [=====] - 1s 21ms/step - loss: 0.3648 -
binary_accuracy: 0.8268 - val_loss: 0.4630 - val_binary_accuracy:
0.8048

Epoch 40/100
51/51 [=====] - 1s 21ms/step - loss: 0.3588 -
binary_accuracy: 0.8341 - val_loss: 0.4675 - val_binary_accuracy:
0.7983
Epoch 41/100
51/51 [=====] - 1s 17ms/step - loss: 0.3657 -
binary_accuracy: 0.8281 - val_loss: 0.4570 - val_binary_accuracy:
0.8014
Epoch 42/100
51/51 [=====] - 1s 15ms/step - loss: 0.3589 -
binary_accuracy: 0.8327 - val_loss: 0.4610 - val_binary_accuracy:
0.8061
Epoch 43/100
51/51 [=====] - 1s 21ms/step - loss: 0.3554 -
binary_accuracy: 0.8348 - val_loss: 0.4743 - val_binary_accuracy:
0.8014
Epoch 44/100
51/51 [=====] - 1s 27ms/step - loss: 0.3567 -
binary_accuracy: 0.8334 - val_loss: 0.4728 - val_binary_accuracy:
0.8033
Epoch 45/100
51/51 [=====] - 1s 17ms/step - loss: 0.3512 -
binary_accuracy: 0.8335 - val_loss: 0.4707 - val_binary_accuracy:
0.8002
Epoch 46/100
51/51 [=====] - 1s 19ms/step - loss: 0.3525 -
binary_accuracy: 0.8353 - val_loss: 0.4581 - val_binary_accuracy:
0.8086
Epoch 47/100
51/51 [=====] - 1s 22ms/step - loss: 0.3504 -
binary_accuracy: 0.8378 - val_loss: 0.4679 - val_binary_accuracy:
0.8123
Epoch 48/100
51/51 [=====] - 1s 16ms/step - loss: 0.3463 -
binary_accuracy: 0.8376 - val_loss: 0.4731 - val_binary_accuracy:
0.8073
Epoch 49/100
51/51 [=====] - 1s 17ms/step - loss: 0.3518 -
binary_accuracy: 0.8354 - val_loss: 0.4664 - val_binary_accuracy:
0.8129
Epoch 50/100
51/51 [=====] - 1s 25ms/step - loss: 0.3402 -
binary_accuracy: 0.8398 - val_loss: 0.4580 - val_binary_accuracy:
0.8173
Epoch 51/100
51/51 [=====] - 1s 26ms/step - loss: 0.3408 -
binary_accuracy: 0.8419 - val_loss: 0.4532 - val_binary_accuracy:
0.8204
Epoch 52/100
51/51 [=====] - 1s 24ms/step - loss: 0.3413 -

binary_accuracy: 0.8400 - val_loss: 0.4548 - val_binary_accuracy:
0.8238
Epoch 53/100
51/51 [=====] - 1s 24ms/step - loss: 0.3368 -
binary_accuracy: 0.8466 - val_loss: 0.4712 - val_binary_accuracy:
0.8139
Epoch 54/100
51/51 [=====] - 1s 25ms/step - loss: 0.3348 -
binary_accuracy: 0.8459 - val_loss: 0.4614 - val_binary_accuracy:
0.8098
Epoch 55/100
51/51 [=====] - 1s 22ms/step - loss: 0.3315 -
binary_accuracy: 0.8480 - val_loss: 0.4558 - val_binary_accuracy:
0.8204
Epoch 56/100
51/51 [=====] - 1s 19ms/step - loss: 0.3277 -
binary_accuracy: 0.8487 - val_loss: 0.4618 - val_binary_accuracy:
0.8139
Epoch 57/100
51/51 [=====] - 1s 12ms/step - loss: 0.3249 -
binary_accuracy: 0.8529 - val_loss: 0.4450 - val_binary_accuracy:
0.8244
Epoch 58/100
51/51 [=====] - 1s 12ms/step - loss: 0.3213 -
binary_accuracy: 0.8515 - val_loss: 0.4637 - val_binary_accuracy:
0.8191
Epoch 59/100
51/51 [=====] - 1s 13ms/step - loss: 0.3269 -
binary_accuracy: 0.8493 - val_loss: 0.4624 - val_binary_accuracy:
0.8266
Epoch 60/100
51/51 [=====] - 1s 12ms/step - loss: 0.3141 -
binary_accuracy: 0.8557 - val_loss: 0.4585 - val_binary_accuracy:
0.8266
Epoch 61/100
51/51 [=====] - 1s 12ms/step - loss: 0.3224 -
binary_accuracy: 0.8543 - val_loss: 0.4775 - val_binary_accuracy:
0.8226
Epoch 62/100
51/51 [=====] - 1s 13ms/step - loss: 0.3195 -
binary_accuracy: 0.8564 - val_loss: 0.4435 - val_binary_accuracy:
0.8247
Epoch 63/100
51/51 [=====] - 1s 13ms/step - loss: 0.3206 -
binary_accuracy: 0.8536 - val_loss: 0.4523 - val_binary_accuracy:
0.8310
Epoch 64/100
51/51 [=====] - 1s 20ms/step - loss: 0.3185 -
binary_accuracy: 0.8551 - val_loss: 0.4626 - val_binary_accuracy:
0.8232

Epoch 65/100
51/51 [=====] - 1s 29ms/step - loss: 0.3177 -
binary_accuracy: 0.8535 - val_loss: 0.4614 - val_binary_accuracy:
0.8241

Epoch 66/100
51/51 [=====] - 1s 27ms/step - loss: 0.3107 -
binary_accuracy: 0.8546 - val_loss: 0.4572 - val_binary_accuracy:
0.8250

Epoch 67/100
51/51 [=====] - 1s 17ms/step - loss: 0.3166 -
binary_accuracy: 0.8596 - val_loss: 0.4464 - val_binary_accuracy:
0.8235

Epoch 68/100
51/51 [=====] - 1s 12ms/step - loss: 0.3145 -
binary_accuracy: 0.8579 - val_loss: 0.4441 - val_binary_accuracy:
0.8310

Epoch 69/100
51/51 [=====] - 1s 12ms/step - loss: 0.3075 -
binary_accuracy: 0.8604 - val_loss: 0.4414 - val_binary_accuracy:
0.8303

Epoch 70/100
51/51 [=====] - 1s 13ms/step - loss: 0.3054 -
binary_accuracy: 0.8616 - val_loss: 0.4483 - val_binary_accuracy:
0.8356

Epoch 71/100
51/51 [=====] - 1s 13ms/step - loss: 0.3121 -
binary_accuracy: 0.8594 - val_loss: 0.4595 - val_binary_accuracy:
0.8282

Epoch 72/100
51/51 [=====] - 1s 13ms/step - loss: 0.3124 -
binary_accuracy: 0.8599 - val_loss: 0.4484 - val_binary_accuracy:
0.8269

Epoch 73/100
51/51 [=====] - 1s 12ms/step - loss: 0.3036 -
binary_accuracy: 0.8621 - val_loss: 0.4516 - val_binary_accuracy:
0.8319

Epoch 74/100
51/51 [=====] - 1s 13ms/step - loss: 0.2926 -
binary_accuracy: 0.8672 - val_loss: 0.4719 - val_binary_accuracy:
0.8229

Epoch 75/100
51/51 [=====] - 1s 13ms/step - loss: 0.3078 -
binary_accuracy: 0.8589 - val_loss: 0.4390 - val_binary_accuracy:
0.8431


Epoch 76/100
51/51 [=====] - 1s 13ms/step - loss: 0.3013 -
binary_accuracy: 0.8677 - val_loss: 0.4459 - val_binary_accuracy:
0.8344

Epoch 77/100
51/51 [=====] - 1s 12ms/step - loss: 0.2995 -

binary_accuracy: 0.8661 - val_loss: 0.4467 - val_binary_accuracy:
0.8424
Epoch 78/100
51/51 [=====] - 1s 13ms/step - loss: 0.2962 -
binary_accuracy: 0.8682 - val_loss: 0.4447 - val_binary_accuracy:
0.8359
Epoch 79/100
51/51 [=====] - 1s 12ms/step - loss: 0.2970 -
binary_accuracy: 0.8670 - val_loss: 0.4522 - val_binary_accuracy:
0.8372
Epoch 80/100
51/51 [=====] - 1s 12ms/step - loss: 0.2921 -
binary_accuracy: 0.8722 - val_loss: 0.4510 - val_binary_accuracy:
0.8353
Epoch 81/100
51/51 [=====] - 1s 12ms/step - loss: 0.2912 -
binary_accuracy: 0.8703 - val_loss: 0.4361 - val_binary_accuracy:
0.8403
Epoch 82/100
51/51 [=====] - 1s 19ms/step - loss: 0.2828 -
binary_accuracy: 0.8770 - val_loss: 0.4403 - val_binary_accuracy:
0.8375
Epoch 83/100
51/51 [=====] - 1s 19ms/step - loss: 0.2873 -
binary_accuracy: 0.8726 - val_loss: 0.4479 - val_binary_accuracy:
0.8372
Epoch 84/100
51/51 [=====] - 1s 20ms/step - loss: 0.2888 -
binary_accuracy: 0.8745 - val_loss: 0.4321 - val_binary_accuracy:
0.8397
Epoch 85/100
51/51 [=====] - 1s 14ms/step - loss: 0.2878 -
binary_accuracy: 0.8701 - val_loss: 0.4457 - val_binary_accuracy:
0.8390
Epoch 86/100
51/51 [=====] - 1s 13ms/step - loss: 0.2810 -
binary_accuracy: 0.8773 - val_loss: 0.4603 - val_binary_accuracy:
0.8387
Epoch 87/100
51/51 [=====] - 1s 13ms/step - loss: 0.2886 -
binary_accuracy: 0.8702 - val_loss: 0.4653 - val_binary_accuracy:
0.8393
Epoch 88/100
51/51 [=====] - 1s 13ms/step - loss: 0.2878 -
binary_accuracy: 0.8724 - val_loss: 0.4552 - val_binary_accuracy:
0.8440
Epoch 89/100
51/51 [=====] - 1s 13ms/step - loss: 0.2735 -
binary_accuracy: 0.8811 - val_loss: 0.4572 - val_binary_accuracy:
0.8484

```
Epoch 90/100
51/51 [=====] - 1s 13ms/step - loss: 0.2729 -
binary_accuracy: 0.8792 - val_loss: 0.4317 - val_binary_accuracy:
0.8449
Epoch 91/100
51/51 [=====] - 1s 13ms/step - loss: 0.2821 -
binary_accuracy: 0.8734 - val_loss: 0.4666 - val_binary_accuracy:
0.8369
Epoch 92/100
51/51 [=====] - 1s 12ms/step - loss: 0.2735 -
binary_accuracy: 0.8794 - val_loss: 0.4583 - val_binary_accuracy:
0.8328
Epoch 93/100
51/51 [=====] - 1s 12ms/step - loss: 0.2807 -
binary_accuracy: 0.8748 - val_loss: 0.4468 - val_binary_accuracy:
0.8468
Epoch 94/100
51/51 [=====] - 1s 12ms/step - loss: 0.2718 -
binary_accuracy: 0.8826 - val_loss: 0.4351 - val_binary_accuracy:
0.8508
```

```
plt.figure(figsize=(15,6))
plt.plot(history.history['binary_accuracy'],color='b',label = 'train
dataset',linewidth=5)
plt.plot(history.history['val_binary_accuracy'],color='g',label='testd
ataset',linewidth=5)
plt.legend()
plt.show()
```



got accuracy around 88% on train set and 85% on test set